

CS 564 - Project Part 4

The Minirel HeapFile Manager

Introduction

In this part of the project, you will implement a File Manager for Heap Files that also provides a scan mechanism that will allow you to search a heap file for records that satisfy a search predicate called a *filter*. How are heap files different from the files that the DB layer of Minirel provides? The main difference is that the files at the DB layer are physical (a bunch of pages) and not logical. The HeapFile layer imposes a logical ordering on the pages (via a linked list).

Classes

The Heapfile layer contains three main classes: a **FileHdrPage** class that implements a heap file using a linked list of pages, a **HeapFile** class for inserting and deleting records from a heap file and a **HeapFileScan** class (derived from the HeapFile class) for scanning a heap file (with an optional predicate). The detailed description of each of these classes follows.

The FileHdrPage Class

Each heap file consists of one instance of the FileHdrPage class plus one or more data pages. This header page is different from the DB header page for the underlying physical file provided by the I/O layer.

```
class FileHdrPage {  
    char fileName[MAXNAMESIZE]; // name of file  
    int firstPage; // pageNo of first data page in file  
    int lastPage; // pageNo of last data page in file  
    int pageCount; // number of pages  
    int recCnt; // record count  
};
```

The meaning of most of the fields in this structure should be obvious from their names. lastPage is used to keep track of the last page in the file. The firstPage attribute of the header page points to the first data page. The data pages are instances of the Page class and they are linked together using a linked list (via the nextPage data member).

The FileHdrPage class has NO member functions. The reason for this will be explained below.

Associated with this class are two functions that you are to implement: **createHeapFile()** and **destroyHeapFile()**.

createHeapFile(const string & filename)

This function creates an empty (well, almost empty) heap file. To do this create a db level file by calling db->createfile(). Then, allocate an empty page by invoking bm->allocPage() appropriately. As you know allocPage() will return a pointer to an empty page in the buffer pool along with the page number of the page. Take the Page* pointer returned from allocPage() and cast it to a FileHdrPage*. Using this pointer initialize the values in the header page. Then make a second call to bm->allocPage(). This page will be the first data page of the file. Using the Page* pointer returned, invoke its init() method to initialize the page contents. Finally, store the page number of the data page in firstPage and lastPage attributes of the FileHdrPage.

When you have done all this unpin both pages and mark them as dirty.

TIPS: If you use a pointer, make sure that pointer has been properly initialized. For example, many students write the first few lines of code for this procedure like this:

```
status = db.openFile(fileName, file);  
if (status != OK) {  
    status = db.createFile(fileName);  
    status = bufMgr->allocPage(file, hdrPageNo, newPage);  
    ...
```

This is wrong. If status != OK, then the file fileName does not exist, and the variable "file" is not initialized. The command db.createFile(fileName) creates a file named fileName on disk. But you still have to open this file, using db.openFile(fileName, file) again. Or else the variable "file" is still not initialized, and so the next command, bufMgr->allocPage(file, hdrPageNo, newPage) will refer to an uninitialized variable "file", and you will get a segmentation fault somewhere. Remember that all pointers that you use must have been initialized correctly.

destroyHeapFile (const string & filename)

This is easy. Simply call db->destroyFile(). The user is expected to have closed all instances of the file before calling this function.

The HeapFile Class

As discussed above, a heap file consists of one file header page and 1 or more data pages. The HeapFile class provides a set of methods to manipulate heap files including adding and deleting records and scanning all records in a file. Creating an instance of the heapFile class opens the heap file and reads the file header page and the first data page into the buffer pool.

Here is the definition of the HeapFile class:

```
class HeapFile {  
protected:  
    File* filePtr; // underlying DB File object  
    FileHdrPage* headerPage; // pinned file header page in buffer pool  
    int headerPageNo; // page number of header page  
    bool hdrDirtyFlag; // true if header page has been updated  
  
    Page* curPage; // data page currently pinned in buffer pool  
    int curPageNo; // page number of pinned page  
    bool curDirtyFlag; // true if page has been updated  
    RID curRec; // rid of last record returned  
  
public:  
  
    // initialize  
    HeapFile(const string & name, Status& returnStatus);  
  
    // destructor  
    ~HeapFile();  
  
    // return number of records in file  
    const int getRecCnt() const;  
  
    // given a RID, read record from file, returning pointer and length  
    const Status getRecord(const RID& rid, Record & rec);  
};
```

The public member functions are described below.

HeapFile(const string & name, Status& returnStatus)

This method first opens the appropriate file by calling db.openFile() (do not forget to save the File* returned in the filePtr data member). Next, it reads and pins the header page for the file in the buffer pool, initializing the private data members headerPage, headerPageNo, and hdrDirtyFlag. You might be wondering how you get the page number of the header page. This is what file->getFirstPage() is used for (see description of the I/O layer)! Finally, read and pin the first page of the file into the buffer pool, initializing the values of curPage, curPageNo, and curDirtyFlag appropriately. Set curRec to NULLRID.

-HeapFile()

The destructor first unpins the header page and currently pinned data page and then calls db.closeFile.

const int getRecCnt() const

Returns the number of records currently in the file (as found in the header page for the file).

const Status getRecord (const RID& rid, Record & rec);

This method returns a record (via the rec structure) given the RID of the record. The private data members curPage and curPageNo should be used to keep track of the current data page pinned in the buffer pool. If the desired record is on the currently pinned page, simply invoke curPage->getRecord(rid, rec) to get the record. Otherwise, you need to unpin the currently pinned page (assuming a page is pinned) and use the pageNo field of the RID to read the page into the buffer pool.

TIPS: In heapfile.C the comment above this method says it "retrieves an arbitrary record from a file". This should be understood as retrieving an arbitrary record given the RID of the record. This does not mean just take a random record from the file and return that.

Further, you should check if curPage is NULL. If yes, you need to read the right page (the one with the requested record on it) into the buffer. Make sure when you do this that you DO THE BOOKKEEPING. That is, you need to set the fields curPage, curPageNo, curDirtyFlag, and curRec of the HeapFile object appropriately. Do not forget to do bookkeeping, or else your code won't work properly.

The HeapFileScan Class

The HeapFile class primarily deals with pages. To insert records into a file, the InsertFileScan class is used (described below). To retrieve records the HeapFileScan class (which also represents a HeapFile since it is derived from the HeapFile class) is used. The HeapFileScan class can be used in three ways:

1. To retrieve all records from a HeapFile.
2. To retrieve only those records that match a specified predicate.
3. To delete records in a file

Several HeapFileScans may be instantiated on the same file simultaneously. This will work fine as long as each has its own "current" FileHdrPage pinned in the buffer pool. Note that the HeapFileScan class is derived from the HeapFile class, thus inheriting its data members and member functions.

```
class HeapFileScan : public HeapFile {  
public:  
  
    HeapFileScan(const string & name, Status & status);  
  
    // end filtered scan  
    ~HeapFileScan();  
  
    const Status startScan(const int offset,  
                          const int length,  
                          const Datatype type,  
                          const char* filter,  
                          const Operator op);  
  
    const Status endScan(); // terminate the scan  
    const Status markScan(); // save current position of scan  
    const Status resetScan(); // reset scan to last marked location  
  
    // return RID of next record that satisfies the scan  
    const Status scanNext(RID& outRid);  
  
    // read current record, returning pointer and length  
    const Status getRecord(Record & rec);  
  
    // delete current record  
    const Status deleteRecord();  
  
    // marks current page of scan dirty  
    const Status markDirty();  
  
private:  
    int offset; // byte offset of filter attribute  
    int length; // length of filter attribute  
    Datatype type; // datatype of filter attribute  
    const char* filter; // comparison value of filter  
    Operator op; // comparison operator of filter  
  
    // The following variables are used to preserve the state  
    // of the scan when the method markScan() is invoked.  
    // A subsequent invocation of resetScan() will cause the  
    // scan to be rolled back to the following  
    int markedPageNo; // value of curPageNo when scan was "marked"  
    RID markedRec; // value of curRec when scan was "marked"  
    const bool matchRec(const Record & rec) const;  
};
```

Private Data Members of HeapFileScan

First note that the protected data members curPage, curPageNo, curRec, and curDirtyFlag should be used to implement the scan mechanism.

The other private data members are used to implement conditional or filtered scans. length is the size of the field on which predicate is applied and offset is its position within the record (we consider fixed length attributes only). The type of the attribute can be STRING, INTEGER or FLOAT and is defined in heapfile.h. Similarly all ordinary comparison operators (as defined in heapfile.h) must be supported. The value to be compared against is stored in binary form and is pointed to by filter.

HeapFileScan(const string& name, Status& status)

Initializes the data members of the object and then opens the appropriate heapfile by calling the HeapFile constructor with the name of the file. The status of all these operations is indicated via the status parameter.

For those of you new to C++, since HeapFileScan is derived from HeapFile, the constructor for the HeapFile class is invoked before the HeapFileScan constructor is invoked.

-HeapFileScan()

Shuts down the scan by calling endScan(). After the HeapFileScan destructor is invoked, the HeapFile destructor will be automatically invoked

const Status startScan(const int offset, const int length, const Datatype type, const char* filter, const Operator op);

This method initiates a scan over a file. If filter == NULL, an unconditional scan is performed meaning that the scan will return all the records in the file. Otherwise, the data members of the HeapFileScan object are initialized with the parameters to the method.

const Status endScan();

This method terminates a scan over a file but does not delete the scan object. This will allow the scan object to be reused for another scan.

const Status markScan();

Saves the current position of the scan by preserving the values of curPageNo and curRec in the private data members markedPageNo and markedRec, respectively.

const Status resetScan();

Resets the position of the scan to the position when the scan was last marked by restoring the values of curPageNo and curRec from markedPageNo and markedRec, respectively. Unless the page number of the currently pinned page is the same as the marked page number, unpin the currently pinned page, then read markedPageNo from disk and set curPageNo, curPage, curRec, and curDirtyFlag appropriately.

const Status scanNext(RID& outRid)

Returns (via the outRid parameter) the RID of the next record that satisfies the scan predicate. The basic idea is to scan the file one page at a time. For each page, use the firstRecord() and nextRecord() methods of the Page class to get the rid of all the records on the page. Convert the rid to a pointer to the record data and invoke matchRec() to determine if record satisfies the filter associated with the scan. If so, store the rid in curRec and return curRec. To make things fast, keep the current page pinned until all the records on the page have been processed. Then continue with the next page in the file. Since the HeapFileScan class is derived from the HeapFile class it also has all the methods of the HeapFile class as well. Returns OK if no errors occurred. Otherwise, return the error code of the first error that occurred.

TIPS: It is always a good idea to check if curPage is NULL, and if so, then how to handle that case.

const Status getRecord(Record& rec)

Returns (via the rec parameter) the record whose rid is stored in curRec. The page containing the record should already be pinned in the buffer pool (by a preceding scanNext() call). If not, return BADPAGENO. A pointer to the pinned page can be found in curPage. Just invoke Page::getRecord() on this page and return its status value.

const Status deleteRecord()

Deletes the record with RID whose rid is stored in curRec from the file by calling Page::deleteRecord. The record must be a record on the "current" page of the scan. Returns BADPAGENO if the page number of the record (i.e. curRec.pageNo) does not match curPageNo, otherwise OK.

const bool matchRec(const Record& rec) const

This private method determines whether the record rec satisfies the predicate associated with the scan. It takes care of attributes not being aligned properly. Returns true if the record satisfies the predicate, false otherwise. We will provide this method for you.

const Status markDirty()

Marks the current page of the scan dirty by setting the dirtyFlag. The dirtyFlag should be reset every time a page is pinned. Make sure that when you unpin a page in the buffer pool, that you pass dirtyFlag as a parameter to the unpin call. Returns OK if no errors.

const Status resetScan()

Saves the current position of the scan by preserving the values of curPageNo and curRec in the private data members markedPageNo and markedRec, respectively.

const Status scanNext(RID& outRid)

Returns (via the outRid parameter) the RID of the next record that satisfies the scan predicate. The basic idea is to scan the file one page at a time. For each page, use the firstRecord() and nextRecord() methods of the Page class to get the rid of all the records on the page. Convert the rid to a pointer to the record data and invoke matchRec() to determine if record satisfies the filter associated with the scan. If so, store the rid in curRec and return curRec. To make things fast, keep the current page pinned until all the records on the page have been processed. Then continue with the next page in the file. Since the HeapFileScan class is derived from the HeapFile class it also has all the methods of the HeapFile class as well. Returns OK if no errors occurred. Otherwise, return the error code of the first error that occurred.

TIPS: If curPage is NULL, then the last page of the current page and read it into the buffer. Call curPage->insertRecord to insert the record. If successful, remember to DO THE BOOKKEEPING. That is, you need to set the fields curPage, curPageNo, curDirtyFlag, and curRec of the HeapFile object appropriately. Do not forget to do bookkeeping, or else your code won't work properly.

Getting Started

Start by downloading the Zip file for this stage (your instructor will give the precise instruction). In the Zip file you will find the following files:

```
Makefile - A make file.  
buf.h - Class definitions for the buffer manager. You should not change this.  
bufC - Implementation of the buffer manager. You should not change this.  
db.h - Class definitions for the I/O layer. You should not change this.  
db.C - Implementation of the I/O layer. You should not change this.  
error.h - Error codes.  
error.C - Error class implementation.  
page.h - Definitions for the Page class. You should not change this.  
heapfile.h - Definitions for the HeapFile and HeapFileScan classes.  
heapfile.C - Skeleton implementations for the above classes. You should complete these.  
testfile.C - Test driver for the HeapFile and HeapFileScan classes. You can augment this.
```

Error Checking

This is some additional clarification about how error checking should be done. Each function that you implement must do some error checking on its input parameters to make sure that they are valid. If they are not, then the function should return an appropriate error code. Each function should also check the return status of every function it calls. If an error occurred, you should just return this code to the higher level calling function. Remember to do any necessary cleanup while doing a premature exit from a function. Usually, you should not print out error messages in any of the functions you implement. These should only be printed by the top level calling function in testfile.C for example. The intermediate functions in the HeapFile and Buffer Manager layers should just relay the error codes to the higher layers.